

A close-up photograph of a snail with a brown, spiraled shell resting on a black computer keyboard. The snail's body is extended, and its two eye stalks are visible. The background is softly blurred, showing more of the keyboard and a hint of a desk surface.

Assignment: Fabian Pascal

Instructions

- This assignment is an individual assignment.
 - Submit your answers by **Friday, 23 February, 2023, 17:00** to Canvas.
 - There is strictly no possibility of late submission.
 - Download the following files from Canvas
“Files > Assignments > Fabian Pascal > Code”.
 - “answers.sql” (template file),
 - “FPTest.sql”,
 - “FPPopulate.sql”, and
 - “FPSchema.sql”.
 - Submit the three queries using the template file “answers.sql” to Canvas
“Assignments > Fabian Pascal>”
Indicate your student number in the space reserved for this purpose in the template file.
-

In 1988, Fabian Pascal, a database designer and programmer and prolific blogger on database issues, (see <http://www.dbdebunk.com>) published the article “SQL Redundancy and DBMS Performance” in the journal Database Programming & Design. He compared and discussed the plan and performance of seven equivalent SQL queries with different database management systems. For the experiment he proposed a schema and a synthetic instance on which the seven queries are executed.

At the time, the different systems could or could not execute all the queries and the performances significantly differed among and within individual systems while one would expect the DBMS optimiser to choose the same optimal execution plan for these queries.

In this project, we propose to replay Fabian Pascal’s experiment with PostgreSQL current version.

1. (0 points) Fabian Pascal proposed a very simple schema with two tables: **employee** and **payroll**. The table **employee** records information about employees of a fictitious company. Employees have an employee identifier, a first name and a last name, an address recorded as a street address, a city, a state and a zip code. The table **payroll** records, for each employee, her bonus and salary.

- (a) Create a database in PostgreSQL. Use the “FPSchema.sql” SQL script to create the tables **employee** and **payroll** with the domains suggested in Fabian Pascal’s original article.

```
1 CREATE TABLE employee (
2   empid CHAR(9),
3   lname CHAR(15),
4   fname CHAR(12),
5   address CHAR(20),
6   city CHAR(20),
7   state CHAR(2),
8   zip CHAR(5));
```

```
1 CREATE TABLE payroll (
2   empid CHAR(9),
3   bonus INTEGER,
4   salary INTEGER);
```

- (b) Populate the database. PL/pgSQL is a procedural language to write code that can be executed by the PostgreSQL server directly. Use the “FPPopulate.sql” SQL script to generate a random instance of the database.

Use the PL/pgSQL “**random_string()**” function to generate random strings of upper case alphabetical characters of a fixed length.

```
1 CREATE or REPLACE FUNCTION random_string(length INTEGER) RETURNS TEXT AS
2 $$
3 DECLARE
4   chars TEXT[] := '{A,B,C,D,E,F,G,H,I,J,K,L,M,N,O,P,Q,R,S,T,U,V,W,X,Y,Z}';
5   result TEXT := '';
6   i INTEGER := 0;
7 BEGIN
8   IF length < 0 then
9     RAISE EXCEPTION 'Given length cannot be less than 0';
10  END IF;
11  FOR i IN 1..length
12  LOOP
13    result := result || chars[1+random()*(array_length(chars, 1)-1)];
14  END LOOP;
15  RETURN result;
16 END;
17 $$ LANGUAGE plpgsql;
```

Use the following SQL DML code to call the function and insert data into the two tables.

```
1 INSERT INTO employee
2   SELECT
3     TO_CHAR(g, '09999') AS empid,
4     random_string(15) AS lname,
5     random_string(12) AS fname,
6     '500 ORACLE PARKWAY' AS address,
7     'REDWOOD SHORES' AS city,
8     'CA' AS state,
9     '94065' AS zip
10  FROM
11    generate_series(0, 9999) g;
```

```

1 INSERT INTO payroll(empid, bonus, salary)
2   SELECT
3     per.empid,
4     0 as bonus,
5     99170 + ROUND(random() * 1000)*100 AS salary
6   FROM
7     employee per;

```

- (c) To measure the planning and execution times of a query, we create a PL/pgSQL function called `test` that takes an SQL query `Q` and a number `N` as its parameters and returns the **average planning and execution times** in milliseconds, as reported by `EXPLAIN ANALYZE Q` over `N` executions of the query `Q`.

The code of the function is given below and is available in the “FPTest.sql” SQL script.

```

1 CREATE OR REPLACE FUNCTION test (TEXT, INT) RETURNS TEXT AS
2 $$
3 DECLARE
4   r RECORD;
5   p TEXT;
6   e TEXT;
7   ap NUMERIC := 0;
8   ae NUMERIC := 0;
9 BEGIN
10  FOR i IN 1..$2
11  LOOP
12    FOR r in EXECUTE 'EXPLAIN ANALYZE ' || $1
13    LOOP
14      IF r::TEXT LIKE '%Planning%'
15      THEN
16        p := regexp_replace( r::TEXT, '.*Planning (?:T|t)ime: (.*?) ms.*', '\1');
17      END IF;
18      IF r::TEXT LIKE '%Execution%'
19      THEN
20        e := regexp_replace( r::TEXT, '.*Execution (?:T|t)ime: (.*?) ms.*', '\1');
21      END IF;
22    END LOOP;
23    ap := ap + (p::NUMERIC - ap) / i;
24    ae := ae + (e::NUMERIC - ae) / i;
25  END LOOP;
26  RETURN ROUND(ap, 2) || ' : ' || ROUND(ae, 2) ;
27 END;
28 $$ LANGUAGE plpgsql;

```

- (d) Run the PL/pgSQL function above with the Query Tool in pgAdmin 4. You can measure the planning and execution times of a query by running the following SQL query with the Query Tool in pgAdmin 4. The following is an example that prints the average planning and execution times over 1000 executions.

```

1 SELECT test('SELECT per.empid, per.lname
2   FROM employee per, payroll pay
3   WHERE per.empid = pay.empid AND pay.salary = 189170;', 1000);

```

2. (4 points) “There’s more than one way to do it” (Perl motto). Fabian Pascal proposed to evaluate the performance of different equivalent queries finding the identifier and the last name of the employees earning a salary of \$189170. Let us consider the following reference SQL query.

```
1 SELECT per.empid, per.lname
2 FROM employee per, payroll pay
3 WHERE per.empid = pay.empid AND pay.salary = 189170
4 ORDER BY per.empid, per.lname;
```

We are looking for different but equivalent SQL queries that produce the same result as the reference query in the same order for database instances generated with the given scripts.

Queries that do not execute with PostgreSQL or do not produce the correct answer shall receive 0 mark.

- (a) (2 points) Complete the following query to produce the same result as the reference query in the same order for database instances generated with the given scripts. Minimally modify the “FROM” and “WHERE” clauses by replacing all and only the occurrences of “TRUE”.

```
1 SELECT per.empid, per.lname
2 FROM employee per RIGHT OUTER JOIN payroll pay
3     ON TRUE AND pay.salary = 189170
4 WHERE TRUE
5 ORDER BY per.empid, per.lname;
```

- (b) (2 points) Complete the following query to produce the same result as the reference query in the same order for database instances generated with the given scripts. Minimally modify the inner query by replacing all and only the occurrences of “TRUE”.

```
1 SELECT per.empid, per.lname
2 FROM employee per
3 WHERE NOT EXISTS (
4     SELECT TRUE)
5 ORDER BY per.empid, per.lname;
```

3. (6 points) The Long Way (just “Don’t repeat yourself”). We investigate constructions that may prevent PostgreSQL from optimising a query well.

- (a) (6 points) Propose a new query that produces the same result as the reference query in the same order for database instances generated with the given scripts, which is as slow (sum of average planning and execution times) as possible.

Do not modify the schema and the data.

The query should not mention the name of any table more than twice.

Do not use sleep functions or SQL constructs not discussed in the lecture or other devices that may be arbitrated as unnecessary at the discretion of the marking team.

The execution of the query must terminate on your machine, and you should be able to measure and indicate its average planning and average execution time of over 1000 executions (the marking team may ask you for a demonstration).

This question is marked competitively on the speed and then on the interest and originality of the answer. Only the slowest, most interesting, and most original queries receive more than 1 mark. Note that identical answers are correspondingly less original. Queries that do not produce the same result in the same order for database instances generated with the given scripts as the reference query receive 0 mark.

– END OF PAPER –